

Solution to the

Gaming Parlor

Programming Project

The Gaming Parlor - Solution

□ Scenario:

- ❖ Front desk with dice (*resource units*)
- ❖ Groups request (e.g., 5) dice (*They request resources*)
- ❖ Groups must wait, if none available
- ❖ A list of waiting groups... A "condition" variable
- ❖ Dice are returned (*resources are released*)
- ❖ The condition is signaled
- ❖ The group checks and finds it needs to wait some more
- ❖ The group (thread) waits...and goes to the end of the line

□ Problem?

□

The Gaming Parlor - Solution

□ Scenario:

- ❖ Front desk with dice (*resource units*)
- ❖ Groups request (e.g., 5) dice (*They request resources*)
- ❖ Groups must wait, if none available
- ❖ Dice are returned (*resources are released*)
- ❖ A list of waiting groups... A "condition" variable
- ❖ The condition is signalled
- ❖ The group checks and finds it needs to wait some more
- ❖ The group (thread) waits...and goes to the end of the line

□ Problem?

- ❖ Starvation!

The Gaming Parlor - Solution

- *Approach:*
 - ❖ Serve every group "first-come-first-served".

- *Implementation:*
 - ❖ Keep the thread at the front of the line separate
 - ❖ "Leader" - the thread that is at the front of the line
 - ❖ Use 2 condition variables.
 - ❖ "Leader" will have at most one waiting thread
 - ❖ "RestOfLine" will have all other waiting threads

The Threads

```
function Group (numDice: int)  
    var i: int  
    for i = 1 to 5  
        gameParlor.Acquire (numDice)  
        currentThread.Yield ()  
        gameParlor.Release (numDice)  
        currentThread.Yield ()  
    endFor  
endFunction
```

```
thA.Init ("A")  
thA.Fork (Group, 4)  
...
```

The Monitor

```
class GameParlor
  superclass Object
  fields
    monitorLock: Mutex
    leader: Condition
    restOfLine: Condition
    numberDiceAvail: int
    numberOfWaitingGroups: int
  methods
    Init ()
    Acquire (numNeeded: int)
    Release (numReturned: int)
    Print (str: String, count: int)
endClass
```

The Release Method

```
method Release (numReturned: int)
    monitorLock.Lock ()

    -- Return the dice
    numberDiceAvail = numberDiceAvail + numReturned

    -- Print
    self.Print ("releases and adds back", numReturned)

    -- Wakeup the first group in line (if any)
    leader.Signal (&monitorLock)

    monitorLock.Unlock ()
endMethod
```

The Acquire Method

```
method Acquire (numNeeded: int)
    monitorLock.Lock ()
    -- Print
    self.Print ("requests", numNeeded)
    -- Indicate that we are waiting for dice.
    numberOfWaitingGroups = numberOfWaitingGroups + 1
    -- If there is a line, then get into it.
    if numberOfWaitingGroups > 1
        restOfLine.Wait (&monitorLock)
    endIf
    -- Now we're at the head of the line. Wait until
        there are enough dice.
    while numberDiceAvail < numNeeded
        leader.Wait (&monitorLock)
    endWhile
    ...
```


The Acquire Method

...

-- Take our dice.

numberDiceAvail = numberDiceAvail - numNeeded

*-- Now we are no longer waiting; wakeup some other
group and leave.*

numberOfWaitingGroups = numberOfWaitingGroups - 1
restOfLine.Signal (&monitorLock)

-- Print

self.Print ("proceeds with", numNeeded)

monitorLock.Unlock ()

endMethod

CS 333
Introduction to Operating Systems

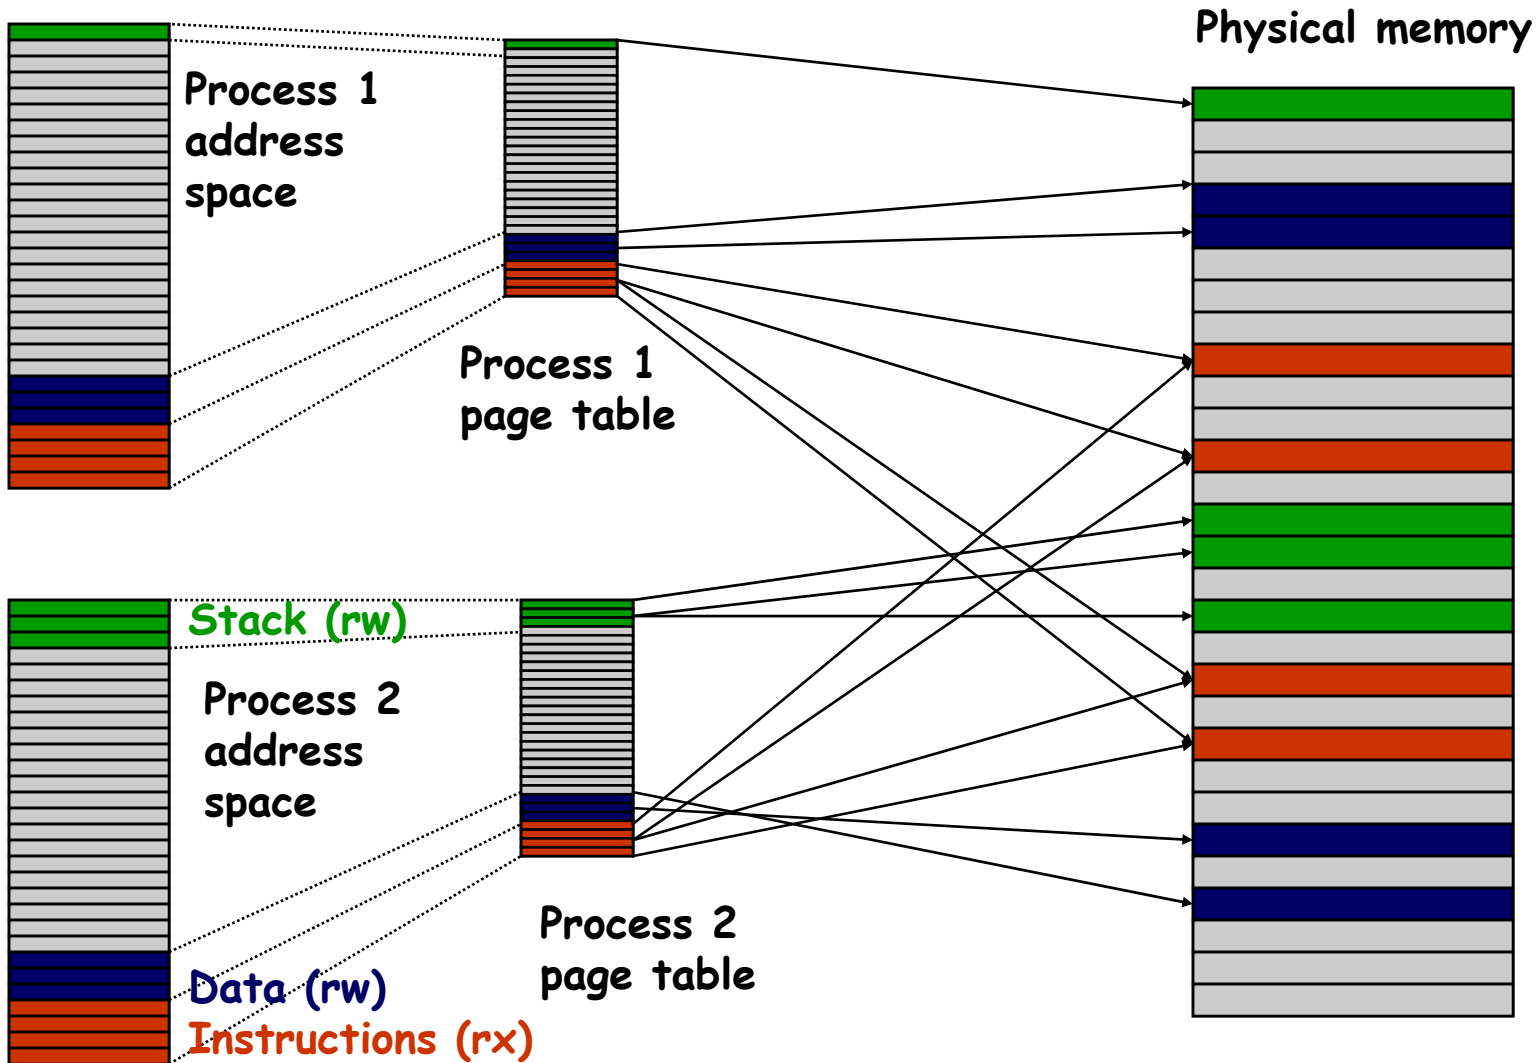
Class 13 - Virtual Memory (3)

Jonathan Walpole
Computer Science
Portland State University

Page sharing

- In a large multiprogramming system...
 - ❖ Some users run the same program at the same time
 - Why have more than one copy of the same page in memory???
- Goal:
 - ❖ Share pages among “processes” (not just threads!)
 - Cannot share writable pages
 - If writable pages were shared processes would notice each other's effects
 - Text segment can be shared

Page sharing



Page sharing

- **“Fork” system call**
 - ❖ Copy the parent's virtual address space
 - ... and immediately do an “Exec” system call
 - Exec overwrites the calling address space with the contents of an executable file (ie a new program)
 - ❖ Desired Semantics:
 - pages are copied, not shared
 - ❖ Observations
 - Copying every page in an address space is expensive!
 - processes can't notice the difference between copying and sharing unless pages are modified!


Page sharing

- **Idea:** **Copy-On-Write**
 - ❖ Initialize new page table, but point entries to existing page frames of parent
 - **Share pages**
 - ❖ Temporarily mark all pages “read-only”
 - **Share all pages until a protection fault occurs**
 - ❖ Protection fault (copy-on-write fault):
 - **Is this page really read only or is it writable but temporarily protected for copy-on-write?**
 - **If it is writable**
 - **copy the page**
 - **mark both copies “writable”**
 - **resume execution as if no fault occurred**

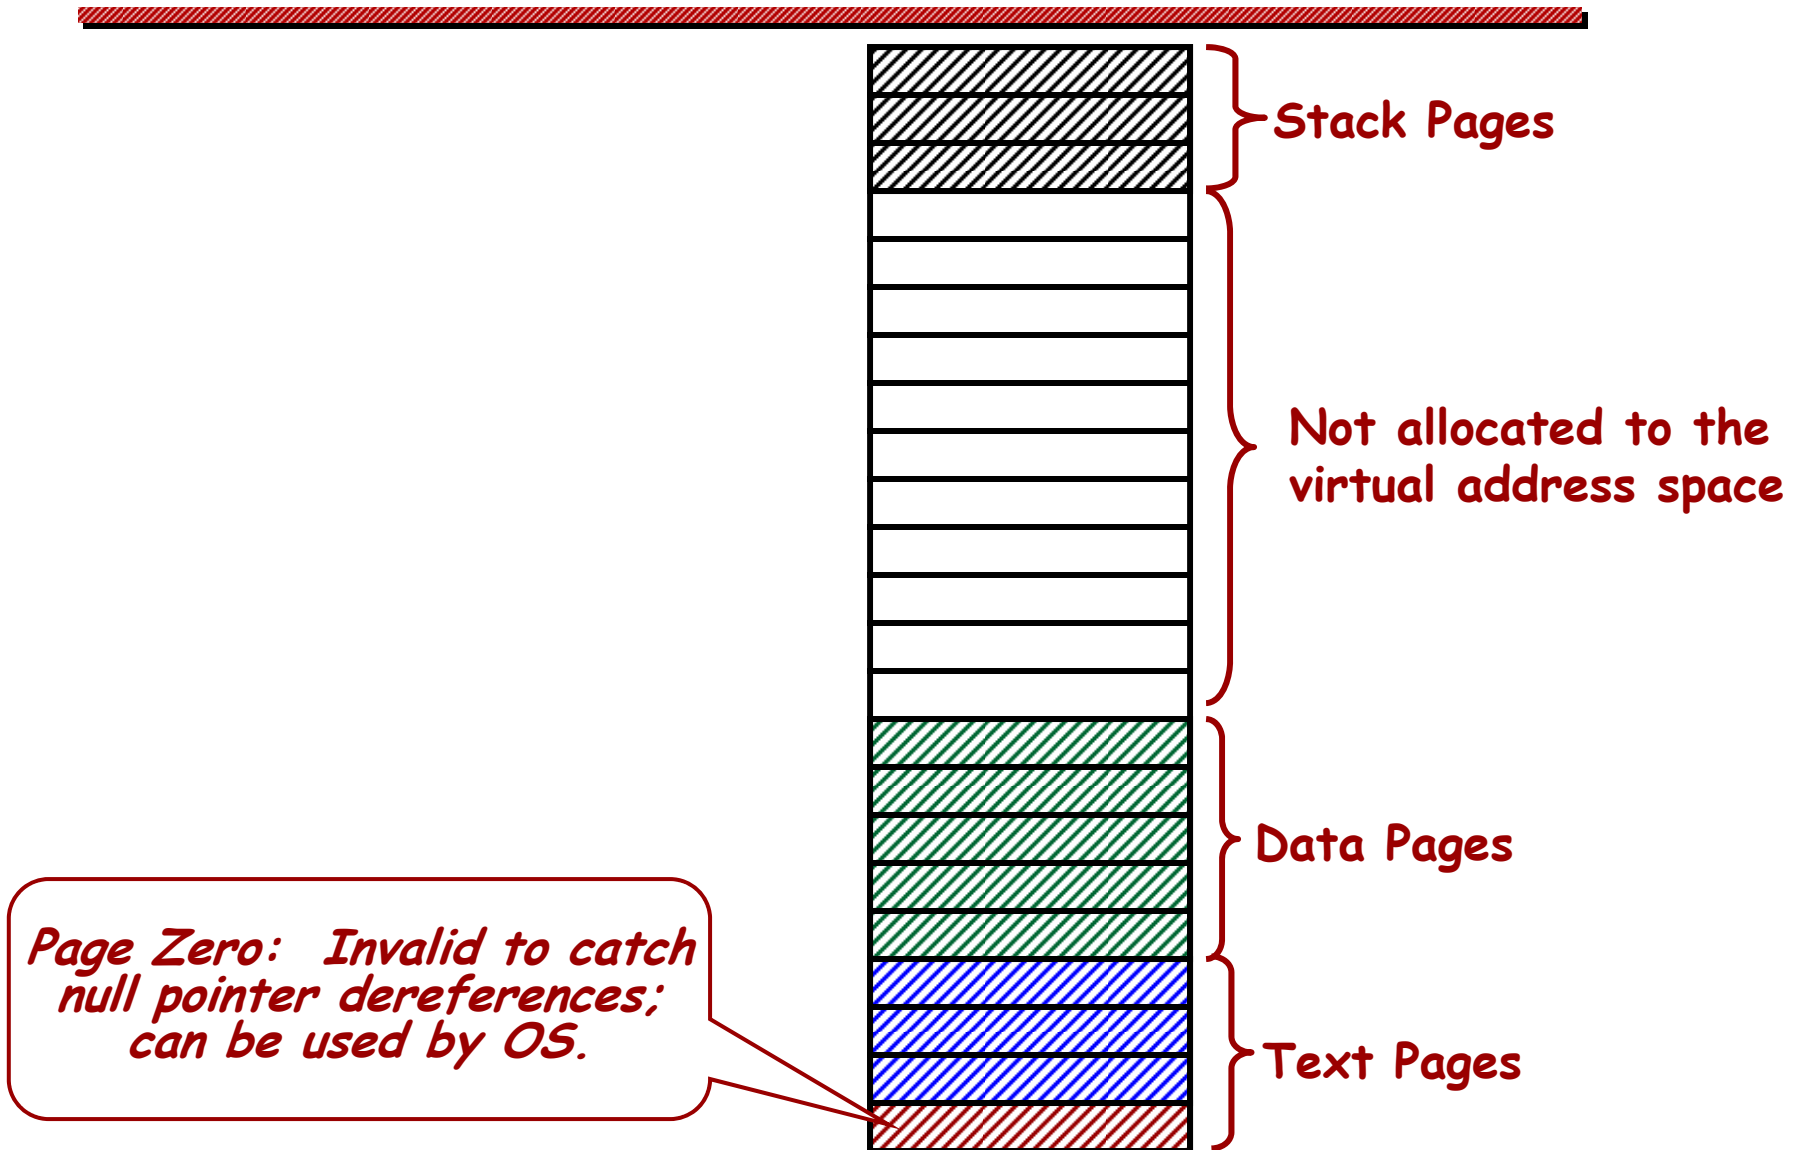
New System Calls for Page Management

- **Goal:**
 - ❖ Allow some processes more control over paging!
- **System calls added to the kernel**
 - ❖ A process can request a page before it is needed
 - Allows processes to grow (heap, stack etc)
 - ❖ Processes can share pages
 - Allows fast communication of data between processes
 - Similar to how threads share memory
 - ... so what is the difference?

Unix processes




Unix processes




Unix processes

The stack grows;
Page requested here




Unix processes

The stack grows;
Page requested here
A new page is allocated
and process continues




Unix processes

The stack grows;
Page requested here
A new page is allocated
and process continues



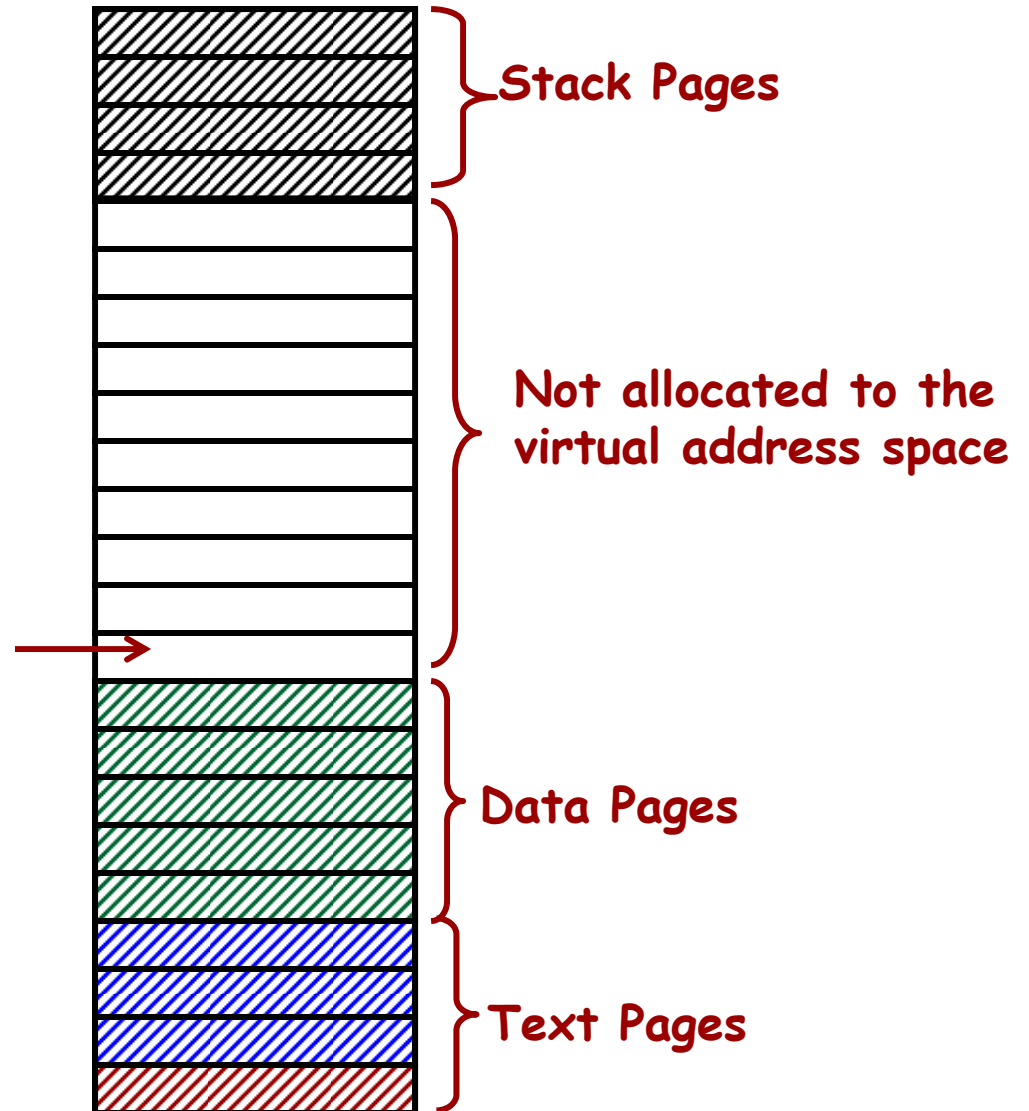
Unix processes

The heap grows;
Page requested here




Unix processes

The heap grows;
Page requested here
A new page is allocated
and process continues



Unix processes

The heap grows;
Page requested here
A new page is allocated
and process continues



Virtual memory implementation

- When is the kernel involved?

Virtual memory implementation

- When is the kernel involved?
 - ❖ *Process creation*
 - ❖ *Process is scheduled to run*
 - ❖ *A fault occurs*
 - ❖ *Process termination*

Virtual memory implementation

- *Process creation*
 - ❖ Determine the process size
 - ❖ Create new page table

Virtual memory implementation

- *Process is scheduled to run*
 - ❖ MMU is initialized to point to new page table
 - ❖ TLB is flushed (unless it's a tagged TLB)

Virtual memory implementation

- *A fault occurs*

- ❖ Could be a TLB-miss fault, segmentation fault, protection fault, copy-on-write fault ...
- ❖ Determine the virtual address causing the problem
- ❖ Determine whether access is allowed, if not terminate the process
- ❖ Refill TLB (TLB-miss fault)
- ❖ Copy page and reset protections (copy-on-write fault)
- ❖ Swap an evicted page out & read in the desired page (page fault)

Virtual memory implementation

- *Process termination*
 - ❖ Release / free all frames (if reference count is zero)
 - ❖ Release / free the page table

Handling a page fault

- ❑ **Hardware traps to kernel**
 - ❖ PC and SR are saved on stack
- ❑ **Save the other registers**
- ❑ **Determine the virtual address causing the problem**
- ❑ **Check validity of the address**
 - ❖ determine which page is needed
 - ❖ may need to kill the process if address is invalid
- ❑ **Find the frame to use (page replacement algorithm)**
- ❑ **Is the page in the target frame dirty?**
 - ❖ If so, write it out (& schedule other processes)
- ❑ **Read in the desired frame from swapping file**
- ❑ **Update the page tables**
- ❑ *(continued)*

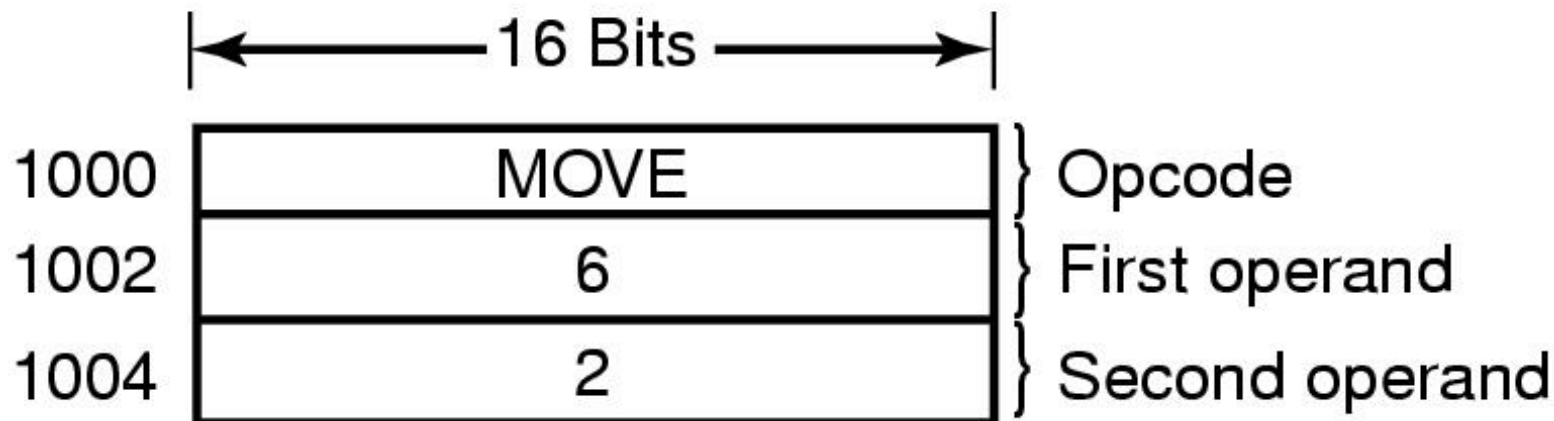
Handling a page fault

- ❑ Back up the current instruction
 - ❖ The "faulting instruction"
- ❑ Schedule the faulting process to run again
- ❑ Return to scheduler
- ❑ ...
- ❑ Reload registers
- ❑ Resume execution

Backing the PC up to restart an instruction

- ❑ Consider a multi-word instruction.
- ❑ The instruction makes several memory accesses.
- ❑ One of them faults.
- ❑ The value of the PC depends on when the fault occurred.
- ❑ How can you know what instruction was executing???

MOVE.L #6(A1), 2(A0)



Solutions

- Lot's of clever code in the kernel
- Hardware support (precise interrupts)
 - ❖ Dump internal CPU state into special registers
 - ❖ Make "hidden" registers accessible to kernel
- What if you swapped out the page containing the first operand in order to bring in the second?


Locking pages in memory

- Virtual memory and I/O interact
 - ❖ Requires "Pinning" pages
- Example:
 - ❖ One process does a read system call
 - (This process suspends during I/O)
 - ❖ Another process runs
 - It has a page fault
 - Some page is selected for eviction
 - The frame selected contains the page involved in the read
- Solution:
 - ❖ Each frame has a flag: "Do not evict me".
 - ❖ Must always remember to un-pin the page!

Managing the swap area on disk

- Approach #1:
 - ❖ A process starts up
 - Assume it has **N pages** in its virtual address space
 - ❖ A region of the swap area is set aside for the pages
 - ❖ There are **N pages** in the swap region
 - ❖ The pages are kept in order
 - ❖ For each process, we need to know:
 - Disk address of page 0
 - Number of pages in address space
 - ❖ Each page is either...
 - In a memory frame
 - Stored on disk

Approach #1




Problem

- *What if the virtual address space grows during execution? i.e. more pages are allocated.*
- Approach #2
 - ❖ Store the pages in the swap in a random order.
 - ❖ View the swap file as a collection of free "swap frames".
 - ❖ Need to evict a frame from memory?
 - Find a free "swap frame".
 - Write the page to this place on the disk.
 - Make a note of where the page is.
 - Use the page table entry.
 - Just make sure the valid bit is still zero!
 - ❖ Next time the page is swapped out, it may be written somewhere else.

Approach #2

This picture uses a separate data structure to tell where pages are stored on disk rather than using the page table

Some information, such as protection status, could be stored at segment granularity



Approach #3

- ❑ **Swap to a file**
 - ❖ Each process has its own swap file
 - ❖ File system manages disk layout of files

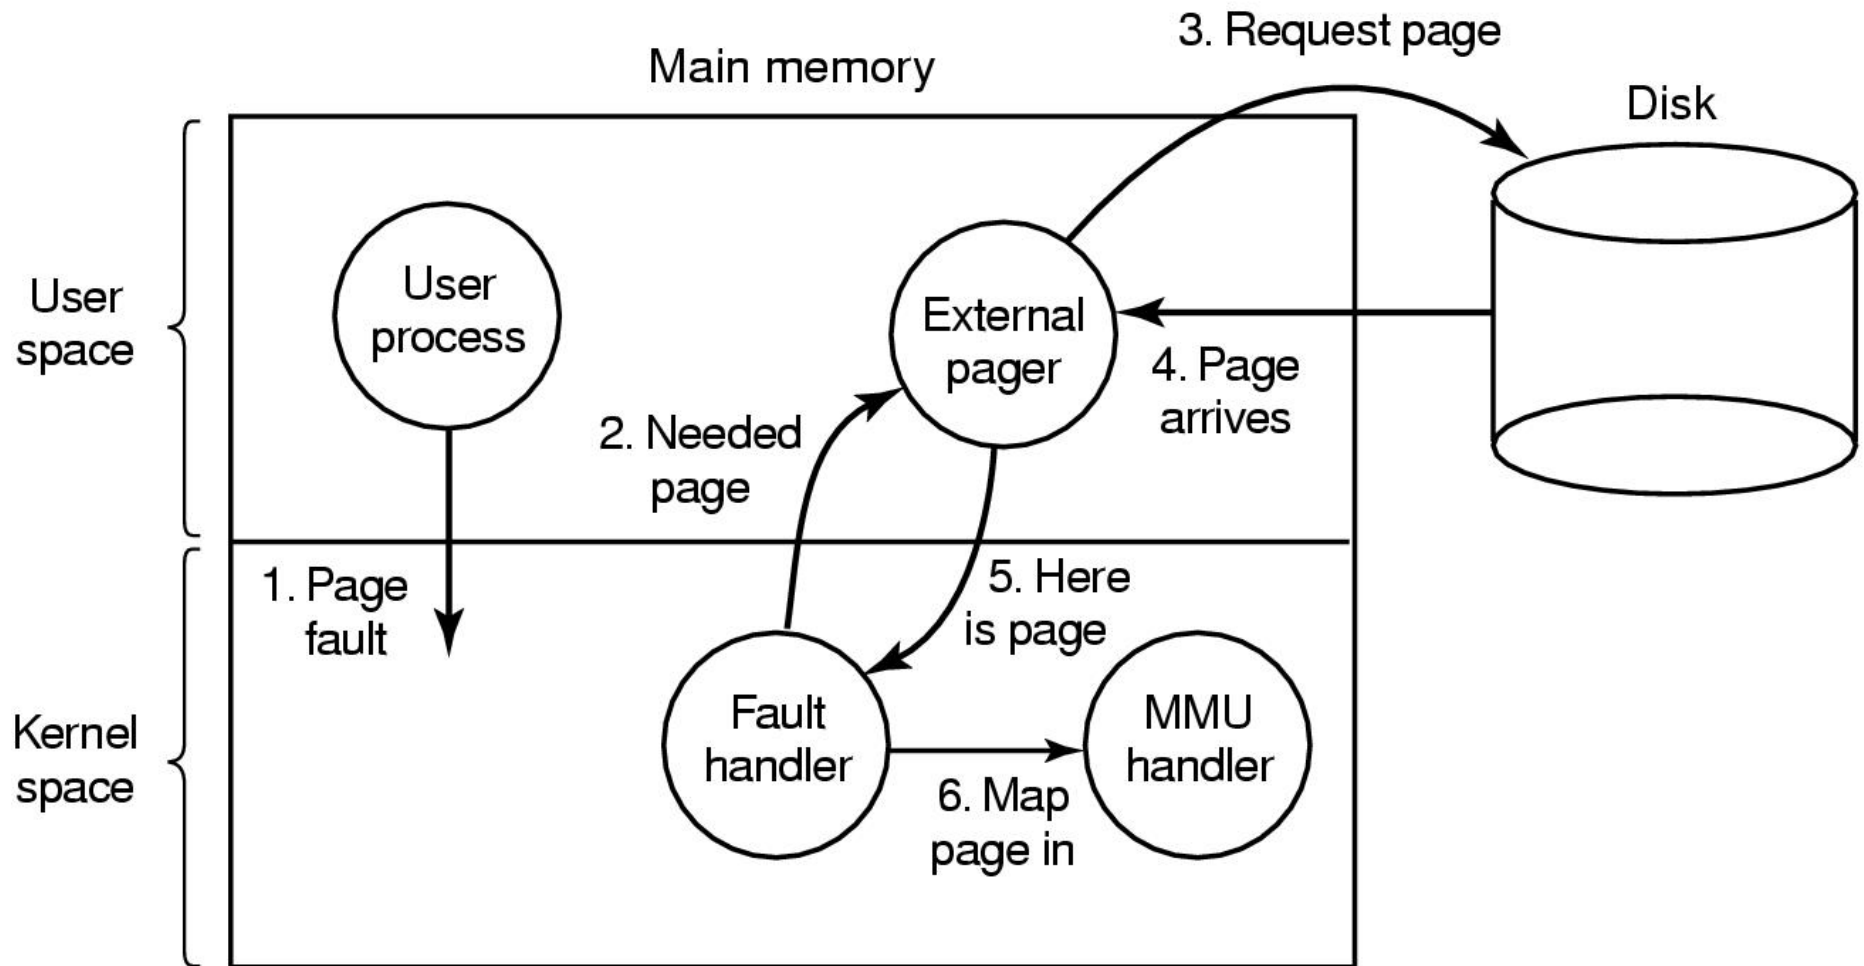
Approach #4

- ❑ Swap to an external pager process (object)
- ❑ A user-level “External Pager” process can determine policy
 - ❖ Which page to evict
 - ❖ When to perform disk I/O
 - ❖ How to manage the swap file
- ❑ When the OS needs to read in or write out a page it sends a message to the external pager
 - ❖ Which may even reside on a different machine
- ❑ Examples: Mach, Minix

Separation of Policy and Mechanism

- Kernel contains
 - ❖ Code to interact with the MMU
 - This code tends to be *machine dependent*
 - ❖ Code to handle page faults
 - This code tends to be *machine independent*

Separation of Policy and Mechanism



Paging performance

- Paging works best if there are plenty of free frames.
- If all pages are full of dirty pages...
 - ❖ Must perform 2 disk operations for each page fault
- It's a good idea to periodically write out dirty pages in order to speed up page fault handling delay

Paging daemon

- **Page Daemon**
 - ❖ A kernel process
 - ❖ Wakes up periodically
 - ❖ Counts the number of free page frames
 - ❖ If too few, run the **page replacement algorithm**...
 - **Select a page & write it to disk**
 - **Mark the page as clean**
 - ❖ If this page is needed later... then it is still there.
 - ❖ If an empty frame is needed later... this page is evicted.